

Jaden Cutinha

Bothell, WA | jadencutinha@gmail.com | 425-318-2881 | jadencutinha.github.io | GitHub | LinkedIn

EDUCATION

Princeton University	Princeton, NJ
<i>B.S.E. Computer Science; Minors in Statistics & Machine Learning and Quantitative Economics</i>	<i>Expected Graduation: May 2028</i>
• Coursework: Data Structures & Algorithms, Programming Systems (C), Computer Architecture, Computational Complexity • Honors: HackPrinceton Fintech Finalist (2026); Princeton Computation & Mathematical Competition Top 5 Finalist (2025) and Honorable Mention (2024); Cerebras AI Research Fellow (2024)	

EXPERIENCE

Sage Investment Group	Jul. 2026 – Present
<i>Private Equity Intern</i>	<i>Kirkland, WA</i>

- Owned a cloud data platform on AWS end to end, from scoping and design through implementation, testing, and deployment, replacing a manual research workflow the team repeated for every target company
- Architected a service-oriented, four-stage pipeline (ingest, normalize, enrich, rank) around a pluggable adapter contract, so a new data source onboards by writing one adapter with no changes to downstream stages
- Enforced schema validation and deduplication gates at every stage boundary to reject malformed or incomplete records before scoring; built an operator-facing dashboard with ranked results and comparative summaries

Sponsors for Educational Opportunities (SEO)	Jun. 2026 – Jul. 2026
<i>Tech Developer Intern</i>	<i>New York, NY</i>

- Implemented a convex-optimization microservice in Python on CVXPY and FastAPI, exposing typed REST endpoints with request validation, backed by asynchronous SQLAlchemy over a SQL database
- Added a Redis cache layer that returns repeated requests without recomputation, avoiding redundant solver runs on identical inputs
- Wrote 250+ pytest unit and integration tests wired into CI/CD with structured logging and health-check endpoints; shipped on a self-organizing Agile Scrum team through sprint planning, Git, and peer code review

BioStrive	Nov. 2025 – Jul. 2026
<i>Software Engineer</i>	<i>Remote</i>

- Built unstructured-document pipelines that extract structured, source-linked fields from raw PDFs with page-level traceability, turning document archives into queryable data products
- Automated validation that checks every extracted value against its source document and flags unsupported values
- Architected multi-tenant Azure infrastructure with per-tenant Cosmos DB document-store isolation, JWT authentication, role-based access control, AES-256 encryption at rest, TLS in transit, and configurable data retention
- Exposed extracted data through versioned REST APIs with input validation, data-quality checks, consistent error handling, and backward-compatible contracts, while keeping onboarding simple for new organizations

The Princeton Laboratory for Artificial Intelligence	Sept. 2025 – May 2026
<i>Artificial Intelligence Research Analyst</i>	<i>Princeton, NJ</i>

- Applied machine learning, LLMs, and NLP to large unstructured text datasets, building sentiment-analysis and pattern-discovery models that fed downstream decision-making; used generative-AI tools to speed up analysis while verifying every output
- Engineered reproducible Python and SQL pipelines to collect, validate, and preprocess unstructured text and time-series data at scale; used held-out validation and statistical testing to guide model selection

Andromeda Labs	May 2025 – Aug. 2025
<i>Machine Learning Engineering Intern</i>	<i>New York, NY</i>

- Built a multi-page full-stack web platform in React and TypeScript with structured data views, filtering, and administrative controls, integrated with backend services through defined API contracts
- Implemented wallet-based authentication and multi-step transaction workflows, validating state at every step so each record stays traceable and reconcilable across frontend, backend, and on-chain systems
- Decoupled frontend and backend so each could evolve and ship independently; turned user and stakeholder needs into shipped features and iterated after launch on usage feedback

PROJECTS

Murmur <i>Rust, Raft, Distributed Systems, Fault Injection, CI</i>
• Built a strongly consistent distributed key/value store from scratch in Rust on a hand-written Raft implementation (no library) • Implemented leader election, log replication, persistence, and fault tolerance, plus a network layer carrying Raft RPCs • Wrote a fault-injection harness that kills nodes under load to verify consistency and recovery, run by CI on every change
Sable <i>Rust, Storage Engines, Crash Recovery, Benchmarking</i>
• Built a durable, crash-recoverable key/value storage engine in Rust with a write-ahead log that logs every write before applying it • Implemented an LSM-tree design with in-memory memtables flushed to immutable SSTables and background compaction • Benchmarked against RocksDB 8.10 on a 200K-key workload: about 50% of its throughput, with negative lookups 2.75x faster

TECHNICAL SKILLS

Languages: Python (pandas, NumPy), Rust, C, SQL, TypeScript, JavaScript, HTML/CSS
Frameworks & Tools: FastAPI, SQLAlchemy, React, Next.js, pytest, Docker, Git, GitHub Actions, CI/CD, REST APIs, OpenAPI
Data & Cloud: AWS, Azure, PostgreSQL, Cosmos DB, Redis, Distributed Systems (Raft), LSM Storage Engines, Data Pipelines/ETL